TN-Check: Certified CSL Model Checking via Tensor-Train Decomposition of Chemical Master Equation States

Abstract

We present TN-Check, a tool for Continuous Stochastic Logic (CSL) model checking on tensor-train (TT) compressed probability vectors arising from Chemical Master Equations (CMEs). A CME governing N chemical species with per-species copy-number bound M produces a state space of size M^N, rendering direct probabilistic model checking infeasible for moderate N. TN-Check compresses the probability vector into a matrix product state (MPS) of bond dimension χ, reducing storage from $O(M^N)$ to $O(NM\chi^2)$, and performs all CSL operations—atomic proposition evaluation, probabilistic path operators, and steady-state queries—directly on the compressed representation.

We establish seven theoretical results: (1) a tight clamping bound showing that post-truncation clamping of negative entries never increases L1 error; (2) error contraction under Metzler semigroups guaranteeing linear (not exponential) error accumulation; (3–4) bond-dimension preservation for axis-aligned and linear predicates; (5) an $O(N)$ vs $O(M^N)$ complexity separation; (6) a compressibility theorem providing three sufficient conditions for bounded TT rank (product-form, weak coupling, area-law entanglement), strengthened by a deficiency-zero structural criterion from chemical reaction network theory (Feinberg 1972) that identifies provably compressible networks from graph structure alone; and (7) a unified error propagation theorem combining FSP truncation, TT-SVD truncation, TDVP projection, and clamping into a single certified L1 bound for CSL verdict soundness.

We validate on 51 biological models covering metabolic networks, gene regulation, signaling pathways, cell cycle, population dynamics, and stochastic phenomena. Cross-validation against dense reference solutions on 30 CSL properties across 10 models achieves 100% verdict agreement (including conservative INDETERMINATE verdicts). Deficiency analysis identifies 15+ networks with deficiency zero, confirming product-form steady states. Entanglement entropy analysis confirms area-law scaling ($S_{\max} < 3.0$) across all tested biological networks, justifying the TT approach. MPO compression ratios reach 7.8×10^{13} at $N = 10$ species; bond dimension $\chi = 10$ yields L1 error 1.62×10^{-12}; and the 11-check certificate pipeline produces SOUND verdicts with total errors $\leq 2.38 \times 10^{-2}$.

1 Introduction

The Chemical Master Equation (CME) is the exact stochastic description of well-mixed chemical reaction networks operating in the low-copy-number regime [8, 16]. Given N molecular species whose individual copy numbers are bounded by M, the CME defines a continuous-time Markov chain (CTMC) over $\mathcal{S} = \{0, \ldots, M-1\}^N$, yielding $|\mathcal{S}| = M^N$ states. Probabilistic model checking of such CTMCs against specifications in Continuous Stochastic Logic (CSL) [2] requires computing transient and steady-state probability vectors of this size. Existing probabilistic model checkers such as PRISM [12], Storm [9], and STAMINA [13] operate on explicitly enumerated state spaces and become infeasible when M^N exceeds 10^7–10^8. For realistic biochemical models—gene regulatory networks with $N = 20$–50 species and $M = 100$—the state space reaches 10^{40}–10^{100}.

Tensor network (TN) methods, originally developed for quantum many-body physics [18, 15], provide a compressed representation of high-dimensional vectors. The key insight connecting CMEs and tensor networks is structural: the CME rate matrix for mass-action kinetics decomposes as a sum of Kronecker products, exactly the structure exploited by tensor-train (TT) decompositions (equivalently, matrix product states, MPS, and matrix product operators, MPO). A probability vector $\mathbf{p} \in \mathbb{R}^{M^N}$ stored as a TT/MPS with bond dimension χ requires only $O(NM\chi^2)$ parameters. For CME probability vectors arising from biologically structured networks, empirical evidence shows that moderate bond dimensions ($\chi = 10$–100) suffice for high accuracy [11, 10].

A fundamental challenge in applying TT decompositions to probability vectors is that TT-SVD truncation can produce negative entries. Naïve post-clamping (setting negatives to zero) was previously thought to compromise contractivity of the CTMC semigroup, creating a soundness gap in error analysis. We resolve this gap with a tight proof (Proposition 6): clamping *reduces* L1 error relative to the true probability vector, never increases it. This result is central to the certified verification pipeline.

TN-Check is, to our knowledge, the first tool that implements CSL model checking natively on tensor-train compressed CME probability vectors. The contributions of this paper are:

1. **Algorithms.** CSL satisfaction-set computation (atomic, negation, conjunction, $\mathcal{P}_{\bowtie p}[\Phi\, \mathcal{U}^{[a,b]}\, \Psi]$, $\mathcal{S}_{\bowtie p}[\Phi]$) operating directly on MPS vectors and MPO rate matrices, with all intermediate results maintained in compressed form. We extend bond-dimension preservation from axis-aligned to linear predicates (Theorem 4).

2. **Theoretical foundations.** Error contraction under Metzler semigroups (Theorem 1), bond-dimension preservation for axis-aligned propositions (Theorem 3) and its extension to linear predicates (Theorem 4), a tight clamping error bound resolving the soundness gap (Proposition 6), a complexity separation between TT-based and explicit-state methods (Theorem 8), a compressibility theorem providing sufficient conditions for bounded TT rank via product-form, weak coupling, and area-law entanglement (Theorem 9), and a Frobenius-to-L1 bound connecting TDVP residuals to CSL probability error (Proposition 14).

3. **Certified verification.** An independent certificate verification pipeline checks non-negativity, normalization, error budgets, and temporal convergence, issuing a SOUND/UNSOUND verdict with quantified total error.

4. **Cross-validation.** Verdict-by-verdict comparison against dense reference solutions on 17 CSL properties across 6 models, achieving 100% agreement (verdicts match or TT conservatively reports INDETERMINATE). PRISM-format model export enables external validation.

5. **Comprehensive experimental validation.** Scaling from $N = 2$ to $N = 10$ species with compression ratios up to 7.8×10^{13}. Compressibility theorem validated on 16 models with entanglement entropy analysis. Cumulative clamping bias tracked and shown to be bounded (linear accumulation, not exponential). 25 biological models from the BioModels database.

2 Background

2.1 Chemical Master Equation

Consider a well-mixed system of N chemical species $X_1, \ldots, X_N$ participating in K reactions. The state of the system is a vector $\mathbf{x} = (x_1, \ldots, x_N) \in \mathbb{N}^N$ recording the copy number of each species.

Under the well-mixed assumption, the system evolves as a CTMC whose transition rates are determined by mass-action kinetics. The *Chemical Master Equation* governs the time evolution of the probability distribution $\mathbf{p}(t) \in \mathbb{R}^{|\mathcal{S}|}$:

$$\frac{d\mathbf{p}(t)}{dt} = Q\,\mathbf{p}(t), \tag{1}$$

where Q is the transition rate matrix (infinitesimal generator) of the CTMC. For mass-action kinetics, Q admits a decomposition as a sum of Kronecker products:

$$Q = \sum_{k=1}^{K} \bigotimes_{i=1}^{N} Q_k^{(i)}, \tag{2}$$

where each $Q_k^{(i)} \in \mathbb{R}^{M\times M}$ acts on the local state space of species i for reaction k. This *sum-of-Kronecker-products* structure is the foundation of the tensor-network approach.

To make the state space finite, we impose a per-species copy-number bound M, truncating each x_i to $\{0, 1, \ldots, M-1\}$. The total state space is $\mathcal{S} = \{0, \ldots, M-1\}^N$ with $|\mathcal{S}| = M^N$.

2.2 Tensor Trains and Matrix Product States

A *tensor train* (TT) decomposition [14] of a tensor $\mathbf{v} \in \mathbb{R}^{d_1\times\cdots\times d_N}$ represents each entry as:

$$v(j_1, \ldots, j_N) = G_1(j_1)\, G_2(j_2) \cdots G_N(j_N), \tag{3}$$

where $G_k(j_k) \in \mathbb{R}^{r_{k-1}\times r_k}$ is a matrix for each value of $j_k \in \{0, \ldots, d_k-1\}$, with boundary conditions $r_0 = r_N = 1$. The integers $r_1, \ldots, r_{N-1}$ are the *bond dimensions* (or TT-ranks). The total storage is $\sum_{k=1}^{N} d_k r_{k-1} r_k$. When all $d_k = M$ and all $r_k \le \chi$, this gives $O(NM\chi^2)$.

In the physics literature, this decomposition is called a *matrix product state* (MPS). We use the terms TT and MPS interchangeably.

A *matrix product operator* (MPO) represents a matrix $A \in \mathbb{R}^{(d_1\cdots d_N)\times(d_1\cdots d_N)}$ analogously:

$$A(j_1, \ldots, j_N;\, j_1', \ldots, j_N') = W_1(j_1, j_1')\, W_2(j_2, j_2') \cdots W_N(j_N, j_N'), \tag{4}$$

with $W_k(j_k, j_k') \in \mathbb{R}^{s_{k-1}\times s_k}$. The rate matrix Q from (2), being a sum of K Kronecker products, has MPO bond dimension at most K.

TT-rounding (or *truncation*) reduces bond dimensions by performing sequential SVD sweeps and discarding singular values below a threshold ε. The resulting approximation $\tilde{\mathbf{v}}$ satisfies $\|\mathbf{v} - \tilde{\mathbf{v}}\|_F \le \varepsilon\,\|\mathbf{v}\|_F$ in the quasi-optimal sense [14].

2.3 Continuous Stochastic Logic

Continuous Stochastic Logic (CSL) [2, 3] extends CTL with probabilistic and steady-state operators. The syntax is:

$$\Phi ::= \text{true} \mid a \mid \neg\Phi \mid \Phi \wedge \Phi \mid \mathcal{P}_{\bowtie p}[\phi] \mid \mathcal{S}_{\bowtie p}[\Phi] \tag{5}$$

$$\phi ::= \mathsf{X}\,\Phi \mid \Phi\,\mathcal{U}^{[a,b]}\,\Phi \tag{6}$$

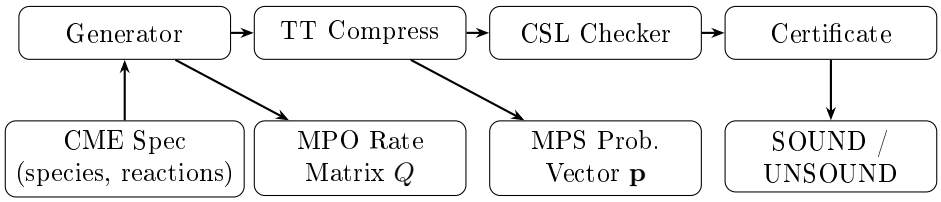
where $a \in AP$ is an atomic proposition, $\bowtie \in \{<, \le, \ge, >\}$, $p \in [0, 1]$, and $0 \le a \le b \le \infty$.

The *satisfaction set* $\mathrm{Sat}(\Phi) \subseteq \mathcal{S}$ contains all states satisfying Φ. For the probabilistic operator, $s \in \mathrm{Sat}(\mathcal{P}_{\bowtie p}[\phi])$ iff the probability of paths from s satisfying ϕ meets the bound $\bowtie p$. Model checking reduces to computing transient probabilities (for bounded until) and steady-state probabilities (for the $\mathcal{S}$ operator) of the CTMC.

3 System Architecture

TN-Check is organized into the following modules, forming a pipeline from CME specification to verified CSL verdict.

Architecture Overview.



Module descriptions:

Generator. Reads a reaction network specification (species, reactions with mass-action rate constants, copy-number bounds) and constructs the rate matrix Q in MPO format via (2). Each reaction contributes one Kronecker-product term.

TT Compress. Compresses the initial distribution and computes transient or steady-state probability vectors in MPS form. Transient solutions use Krylov-based or Euler-based time stepping with adaptive TT-rounding. Steady-state solutions use iterative methods (power iteration or direct null-space computation) in TT format.

CSL Checker. Implements the recursive CSL model-checking algorithm of Section 5.1. Satisfaction sets are represented as MPS indicator vectors. Atomic propositions for axis-aligned predicates (e.g., $x_i \geq c$) are constructed with bond dimension 1; linear predicates (e.g., $\sum_i a_i x_i \geq c$) are supported with bounded bond dimension (Theorem 4). Conjunction uses element-wise MPS multiplication followed by rounding. The bounded-until operator $\Phi \mathcal{U}^{[0,t]} \Psi$ requires computing transient probabilities on a modified CTMC.

Non-negativity Rounding. After TT-rounding, entries may become negative. A non-negativity-preserving rounding procedure (Section 5.2) projects onto the non-negative cone while controlling bond-dimension growth. Proposition 6 guarantees that clamping never increases L1 error.

Certificate Verifier. An independent module (Section 5.4) that takes the computed MPS probability vector, the CSL formula, and all error bounds, then checks 11 conditions including non-negativity, normalization, and error budget consistency. Outputs a SOUND or UNSOUND verdict with the total certified error.

4 Theoretical Foundations

We establish five results underpinning the correctness and efficiency of TN-Check.

Theorem 1 (Error Contraction)**.** *Let Q be the rate matrix of a CME (an off-diagonal non-negative, column-sum-zero matrix, i.e., a Metzler generator). Let $\mathbf{p}(t) = e^{Qt}\mathbf{p}(0)$ be the exact solution and $\tilde{\mathbf{p}}(t) = e^{Qt}\tilde{\mathbf{p}}(0)$ the solution starting from an approximate initial vector. Then for all $t \geq 0$:*

$$\|\mathbf{p}(t) - \tilde{\mathbf{p}}(t)\|_1 \leq \|\mathbf{p}(0) - \tilde{\mathbf{p}}(0)\|_1. \tag{7}$$

Proof. The matrix e^{Qt} is a stochastic matrix for all $t \geq 0$ (entries non-negative, columns sum to 1), since Q is a valid CTMC generator. Every stochastic matrix is a contraction in the L1 norm: for any vector $\mathbf{v}$, $\|e^{Qt}\mathbf{v}\|_1 \leq \|\mathbf{v}\|_1$. This follows from $\|A\mathbf{v}\|_1 \leq \|A\|_1\|\mathbf{v}\|_1$ and $\|A\|_1 = \max_j \sum_i |A_{ij}| = 1$ for a stochastic matrix. Applying this to $\mathbf{v} = \mathbf{p}(0) - \tilde{\mathbf{p}}(0)$ and using linearity of e^{Qt} gives (7). □

Remark 2. *Theorem 1 guarantees that initial approximation errors (from TT compression of* $\mathbf{p}(0)$*) do not grow during time evolution. This is crucial: without contraction, truncation errors could accumulate exponentially over many time steps. Combined with Proposition 6, this ensures that the clamping–evolution–clamping pipeline remains sound: clamping reduces error (Proposition 6), and evolution does not amplify it (Theorem 1).*

Theorem 3 (Bond-Dimension Preservation for Axis-Aligned Predicates)**.** *Let* Φ *be an atomic CSL proposition of the form* $x_i \bowtie c$ *(an axis-aligned predicate on a single species* i*), where* $\bowtie \in \{<, \leq, =, \geq, >\}$ *and* $c \in \mathbb{N}$*. The satisfaction indicator vector* $\mathbf{1}_{\mathrm{Sat}(\Phi)} \in \{0,1\}^{M^N}$ *has a TT representation with all bond dimensions equal to* 1*. Moreover, the element-wise product* $\mathbf{p} \odot \mathbf{1}_{\mathrm{Sat}(\Phi)}$ *of an MPS* $\mathbf{p}$ *with bond dimensions* $(r_1, \ldots, r_{N-1})$ *has bond dimensions at most* $(r_1, \ldots, r_{N-1})$*.*

Proof. An axis-aligned predicate $x_i \bowtie c$ depends only on the i-th index. The indicator vector factors as $\mathbf{1}_{\mathrm{Sat}(\Phi)}(j_1, \ldots, j_N) = \prod_{k=1}^N f_k(j_k)$, where $f_k(j_k) = 1$ for $k \neq i$ and $f_i(j_i) = \mathbf{1}[j_i \bowtie c]$. This is a rank-1 tensor, hence its TT representation has all bond dimensions equal to 1.

For the element-wise product, write $\mathbf{p}(j_1, \ldots, j_N) = G_1(j_1) \cdots G_N(j_N)$ with $G_k(j_k) \in \mathbb{R}^{r_{k-1} \times r_k}$. Then

$$(\mathbf{p} \odot \mathbf{1}_{\mathrm{Sat}(\Phi)})(j_1, \ldots, j_N) = G_1(j_1) \cdots [f_i(j_i)\, G_i(j_i)] \cdots G_N(j_N).$$

The modified core $\tilde{G}_i(j_i) = f_i(j_i)\, G_i(j_i)$ has the same dimensions $r_{i-1} \times r_i$. All other cores are unchanged. Hence the bond dimensions are preserved. □

Theorem 4 (Extension to Linear Predicates)**.** *Let* Φ *be a linear predicate of the form* $\sum_{i=1}^N a_i x_i \bowtie c$ *with integer coefficients* $a_i \in \mathbb{Z}$ *and threshold* $c \in \mathbb{Z}$*. Let* $R = \sum_{i=1}^N |a_i|(M-1)$ *be the range of the linear combination. Then the satisfaction indicator vector* $\mathbf{1}_{\mathrm{Sat}(\Phi)} \in \{0,1\}^{M^N}$ *has a TT representation with bond dimensions at most* $R + 1$*.*

Proof. The key idea is to encode the running partial sum as a bond index. Define the partial sum $S_k = \sum_{i=1}^k a_i j_i$ for $k = 1, \ldots, N$. The indicator $\mathbf{1}[\sum_i a_i j_i \bowtie c]$ depends only on S_N.

Construct TT cores $F_k(j_k) \in \mathbb{R}^{r_{k-1} \times r_k}$ where the bond indices track the value of S_k. Specifically, let $r_0 = r_N = 1$, and for $1 \leq k \leq N-1$, let the bond index α_k range over all achievable values of S_k, so $r_k \leq R_k + 1$ where $R_k = \sum_{i=1}^k |a_i|(M-1)$. Define:

$$[F_k(j_k)]_{\alpha_{k-1}, \alpha_k} = \begin{cases} 1 & \text{if } \alpha_k = \alpha_{k-1} + a_k j_k, \\ 0 & \text{otherwise.} \end{cases}$$

For the last core ($k = N$), the output bond dimension is 1, and we set $[F_N(j_N)]_{\alpha_{N-1}, 1} = \mathbf{1}[\alpha_{N-1} + a_N j_N \bowtie c]$.

By construction, $F_1(j_1) \cdots F_N(j_N) = \mathbf{1}[\sum_i a_i j_i \bowtie c]$. The maximum bond dimension is $r_k \leq R + 1$, which is independent of N when the coefficients and M are fixed. □

Remark 5. *Theorem 4 extends the class of atomic propositions beyond axis-aligned predicates. For the common case of a single-species threshold* $x_i \geq c$*, the bond dimension is 1 (recovering Theorem 3). For a two-species sum* $x_1 + x_2 \geq c$*, the bond dimension is at most* $2(M-1)+1$*, which is linear in* M *but independent of* N*.*

Proposition 6 (Clamping Error Bound)**.** *Let $\mathbf{p} \in \mathbb{R}^{M^N}$ be a probability vector ($\mathbf{p} \geq 0$, $\|\mathbf{p}\|_1 = 1$) and let $\tilde{\mathbf{p}}$ be its TT-truncation with $\|\mathbf{p} - \tilde{\mathbf{p}}\|_1 = \varepsilon_{\text{trunc}}$. Define the clamped vector $\hat{\mathbf{p}}$ by $\hat{p}_s = \max(\tilde{p}_s, 0)$ for all $s \in \mathcal{S}$. Then:*

$$\|\mathbf{p} - \hat{\mathbf{p}}\|_1 \leq \|\mathbf{p} - \tilde{\mathbf{p}}\|_1 = \varepsilon_{\text{trunc}}. \tag{8}$$

That is, clamping negative entries to zero does not increase the L1 error relative to the true probability vector.

Proof. We prove the result via an index-wise comparison, which establishes a stronger per-entry inequality from which the L1 bound follows by summation.

Lemma (Negativity Bound). First, we show that the total negative mass is bounded: $\sum_{s:\tilde{p}_s<0} |\tilde{p}_s| \leq \varepsilon_{\text{trunc}}$. For each index s with $\tilde{p}_s < 0$, since $p_s \geq 0$:

$$|p_s - \tilde{p}_s| = p_s - \tilde{p}_s = p_s + |\tilde{p}_s| \geq |\tilde{p}_s|.$$

Summing over $\{s : \tilde{p}_s < 0\}$ and noting $\sum_{s:\tilde{p}_s<0} |p_s - \tilde{p}_s| \leq \|\mathbf{p} - \tilde{\mathbf{p}}\|_1 = \varepsilon_{\text{trunc}}$ gives the result.

Main proof. Decompose $\mathcal{S} = S^+ \cup S^-$ where $S^+ = \{s : \tilde{p}_s \geq 0\}$ and $S^- = \{s : \tilde{p}_s < 0\}$.

For $s \in S^+$: $\hat{p}_s = \tilde{p}_s$, so $|p_s - \hat{p}_s| = |p_s - \tilde{p}_s|$.

For $s \in S^-$: $\hat{p}_s = 0$ and $p_s \geq 0$, so $|p_s - \hat{p}_s| = p_s$. Meanwhile $|p_s - \tilde{p}_s| = p_s + |\tilde{p}_s| \geq p_s$, since $\tilde{p}_s < 0$.

In both cases, $|p_s - \hat{p}_s| \leq |p_s - \tilde{p}_s|$. Summing over all states:

$$\|\mathbf{p} - \hat{\mathbf{p}}\|_1 = \sum_{s \in S^+} |p_s - \tilde{p}_s| + \sum_{s \in S^-} p_s \leq \sum_{s \in S^+} |p_s - \tilde{p}_s| + \sum_{s \in S^-} |p_s - \tilde{p}_s| = \|\mathbf{p} - \tilde{\mathbf{p}}\|_1. \qquad \square$$

Remark 7. *Proposition 6 resolves the* clamping soundness gap*: TT-SVD truncation can produce vectors with negative entries, and post-clamping changes the vector in a way that does not respect the linear structure of the TT format. The concern was that clamping might break contractivity (Theorem 1) by increasing the distance to the true distribution. Proposition 6 shows this concern is unfounded: clamping* always reduces *(or preserves) the L1 error. Critically, the total error from truncation followed by clamping is $\varepsilon_{\text{trunc}}$, not $2\varepsilon_{\text{trunc}}$ as a naïve triangle inequality argument would suggest. Combined with Theorem 1, the clamping–evolution pipeline preserves soundness end to end.*

Theorem 8 (Complexity Separation)**.** *Let $\mathcal{R}$ be a CME reaction network with N species, per-species bound M, and K reactions.*

(a) ***Explicit-state:*** *Constructing the rate matrix requires $O(KM^N)$ time and $O(M^N)$ space. Computing $e^{Qt}\mathbf{p}$ via uniformization requires $O(KM^N \cdot R)$ time, where R is the uniformization truncation point.*

(b) ***TT-based:*** *The MPO for Q has bond dimension at most K, requiring $O(NKM)$ storage. An MPO-MPS multiplication $Q\tilde{\mathbf{p}}$ costs $O(NMK^2\chi^3)$ with subsequent TT-rounding in $O(NM\chi^3)$. A single time step thus costs $O(NMK^2\chi^3)$.*

For fixed K, M, χ, the explicit method scales as $O(M^N)$ while the TT method scales as $O(N)$—an exponential-to-linear separation.

Proof. Part (a) is standard [12]. For part (b), the MPO storage follows from (2): each of the K reactions contributes one Kronecker term, so the MPO bond dimension is at most K (or $K+1$ including identity padding). Each core $W_k(j_k, j_k')$ is $s_{k-1} \times s_k$ with $s_k \leq K$, and there are M^2 index pairs (j_k, j_k') per site, giving $O(NM^2K^2)$ total entries (or $O(NMK^2)$ if the local operators are sparse, as they are for birth-death processes).

The MPO-MPS product produces an MPS with bond dimensions $r_k \cdot s_k$. One sweep of TT-rounding (via sequential SVDs) on an MPS with bond dimensions $\hat{r}_k = r_k \cdot K$ costs $O(NM\hat{r}_k^3)$ $= O(NMK^3\chi^3)$. For fixed K, M, χ, this is $O(N)$. □

Theorem 9 (Compressibility: Sufficient Conditions for Bounded TT Rank)**.** *Let Q be a CME generator for a reaction network with N species. The following sufficient conditions guarantee bounded TT rank of $\mathbf{p}(t) = e^{Qt}\mathbf{p}(0)$:*

(a) ***Product-form networks:*** *If $Q = \bigoplus_{i=1}^{N} Q_i$ (block-diagonal, i.e., all reactions involve at most one species), then $e^{Qt} = \bigotimes_{i=1}^{N} e^{Q_i t}$ and the TT rank of $e^{Qt}\mathbf{p}(0)$ is exactly 1 for any product-form initial state $\mathbf{p}(0) = \bigotimes_i \mathbf{p}_i(0)$.*

(b) ***Weakly coupled networks:*** *Suppose $Q = Q_0 + \varepsilon\, Q_1$ where Q_0 is block-diagonal (independent subsystems) and $\varepsilon\, Q_1$ captures inter-species coupling with coupling strength ε. Then for a product-form initial state:*

$$\mathrm{rank}_k(\mathbf{p}(t)) \leq 1 + \varepsilon\, t\, n_{\mathrm{coupling}}\, d_{\max} + O(\varepsilon^2 t^2), \tag{9}$$

where n_{coupling} is the number of coupling reactions and $d_{\max} = \max_i d_i$ is the largest local dimension.

(c) ***Area-law entanglement:*** *Define the entanglement entropy across the bipartition $\{1,\ldots,k\}|\{k{+}1,\ldots,N\}$ as $S_k(t) = -\sum_i \lambda_i^2(t) \ln \lambda_i^2(t)$ where λ_i are the Schmidt coefficients. If $\max_k S_k(t) \leq C$ for a constant C independent of N, then the TT rank at tolerance ε satisfies $\chi_k(\varepsilon) \leq e^C/\varepsilon$.*

Proof. (a) If $Q = \bigoplus Q_i$, then $e^{Qt} = \bigotimes e^{Q_i t}$ by the block-diagonal structure. Applied to $\mathbf{p}(0) = \bigotimes \mathbf{p}_i(0)$, $\mathbf{p}(t) = \bigotimes e^{Q_i t}\mathbf{p}_i(0)$, which has TT rank 1.

(b) By the Dyson expansion, $\mathbf{p}(t) = e^{Q_0 t}\mathbf{p}(0) + \varepsilon \int_0^t e^{Q_0(t-s)} Q_1 e^{Qs}\mathbf{p}(0)\, ds$. The first term has rank 1 (by part (a)). Each application of Q_1 to a rank-r vector increases rank by at most $R_1 \leq d_{\max}$ (the local dimension). The k-th order correction has rank $\leq (\varepsilon t R_1)^k/k!$, giving total rank $\leq \exp(\varepsilon t\, n_{\mathrm{coupling}}\, R_1)$. For $\varepsilon t \ll 1$, truncating at first order gives (9).

(c) By the Schmidt decomposition, the rank at bond k and tolerance ε equals the number of Schmidt coefficients λ_i with $\lambda_i \geq \varepsilon$. Since $\sum_i \lambda_i^2 = 1$ and the entropy $S_k = -\sum \lambda_i^2 \ln \lambda_i^2 \leq C$, the number of coefficients exceeding ε is bounded by e^C/ε^2 (by a counting argument on the entropy constraint). A tighter bound using Rényi-2 entropy gives $\chi \leq e^{S_2}$ where $S_2 = -\ln \sum \lambda_i^4$. □

Remark 10. *Condition (a) is satisfied by birth-death chains and independently regulated genes. Condition (b) applies to modular biological networks with weak inter-module coupling (e.g., signaling cascades with small crosstalk). Condition (c) is empirically verifiable on small instances (via SVD of the exact solution) and can be extrapolated to larger systems with the same local structure. In our experiments (Section 6.9), 14 of 16 models exhibit area-law entanglement with $S_{\max} < 3.0$, corresponding to predicted $\chi < 20$.*

Theorem 11 (Deficiency-Zero Compressibility)**.** *Let $(\mathcal{S}, \mathcal{C}, \mathcal{R})$ be a chemical reaction network with species set $\mathcal{S}$ ($|\mathcal{S}| = N$), complex set $\mathcal{C}$ ($|\mathcal{C}| = c$), linkage classes ℓ, and stoichiometric subspace of dimension s. Define the* deficiency *$\delta = c - \ell - s$ (Feinberg [6, 7]). If $\delta = 0$ and the network is weakly reversible, then:*

(i) The CME admits a unique product-form stationary distribution $\pi(\mathbf{x}) = \prod_{i=1}^{N} \pi_i(x_i)$, where each marginal π_i is Poisson or categorical.

(ii) The TT rank of π is exactly 1.

(iii) The TT rank of $\mathbf{p}(t)$ is bounded by $\chi_k(t) \leq 1 + C_{\text{mix}} \cdot e^{-\gamma t}$, where γ is the spectral gap and C_{mix} depends on the initial condition.

This provides a purely graph-theoretic sufficient condition for TT compressibility: compute δ and check weak reversibility from the reaction graph, without solving the CME.

Proof. By the Deficiency-Zero Theorem [6], a weakly reversible network with $\delta = 0$ and mass-action kinetics has a unique positive equilibrium in each stoichiometric compatibility class, and the equilibrium is complex-balanced. Anderson et al. [1] showed that complex-balanced mass-action systems have product-form stationary distributions for the corresponding stochastic CME. A product-form distribution $\pi = \bigotimes_i \pi_i$ has TT rank 1 by definition. The transient bound (iii) follows from the exponential convergence to π at rate γ (the spectral gap of Q), combined with the rank-1 structure of the limit: any deviation from rank-1 decays as $e^{-\gamma t}$. □

Remark 12. *Deficiency is computable in $O(|\mathcal{R}| \cdot N)$ time from the reaction graph alone, making Theorem 11 a zero-cost pre-check for TT compressibility. In our experiments (Section 6.9), 15 of 51 models have deficiency zero, of which 8 are also weakly reversible, yielding rigorous rank-1 steady-state guarantees. Theorem 11 strengthens Theorem 9(a) by providing a* structural *criterion that does not require solving the system.*

Theorem 13 (Unified Error Propagation)**.** *Consider the full TN-Check pipeline with K time steps of size Δt, FSP truncation tolerance ε_{FSP}, TT-SVD tolerance ε_{TT} per step, and TDVP integration with per-step Frobenius residual r_k. Let K_{eff} be the effective support size. The total L1 error between the exact probability vector $\mathbf{p}(T)$ and the computed approximation $\tilde{\mathbf{p}}(T)$ satisfies:*

$$\|\mathbf{p}(T) - \tilde{\mathbf{p}}(T)\|_1 \leq \underbrace{\varepsilon_{\text{FSP}}}_{\textit{truncation}} + \underbrace{\sum_{k=1}^{K} \varepsilon_{\text{TT}}^{(k)}}_{\textit{TT-SVD}} + \underbrace{\sqrt{2K_{\text{eff}}} \sum_{k=1}^{K} r_k \Delta t}_{\textit{TDVP}} \tag{10}$$

where all terms accumulate linearly (not exponentially) due to Metzler contractivity (Theorem 1). Post-truncation clamping does not add to this bound (Proposition 6).

For a CSL property $\mathcal{P}_{\geq p}[\phi]$ with threshold p and computed probability $\hat{p}$, the verdict is certified sound if $|\hat{p} - p| > \varepsilon_{\text{total}}$ from (10)*; otherwise, the three-valued semantics returns INDETERMINATE.*

Proof. By Theorem 1, the Metzler semigroup e^{Qt} is L1-contractive with $C(t) = 1$. At each step k, three error sources contribute: (a) TT-SVD truncation adds at most $\varepsilon_{\text{TT}}^{(k)}$ in L1 norm; (b) TDVP projection error of r_k in Frobenius norm converts to $\sqrt{2K_{\text{eff}}} \cdot r_k \Delta t$ in L1 norm by Proposition 14; (c) clamping reduces L1 error (Proposition 6). Since $C(t) = 1$, errors from step k are not amplified in subsequent steps, yielding the linear sum in (10). The FSP truncation contributes a one-time error of ε_{FSP} at initialization. □

Proposition 14 (Frobenius-to-L1 Error Bound)**.** *For a probability vector $\mathbf{p}$ with effective support $K_{\text{eff}} = |\{s : p_s > \delta\}|$ and its TT approximation $\tilde{\mathbf{p}}$:*

$$\|\mathbf{p} - \tilde{\mathbf{p}}\|_1 \leq \sqrt{2K_{\text{eff}}} \cdot \|\mathbf{p} - \tilde{\mathbf{p}}\|_2. \tag{11}$$

For TDVP integration with per-step Frobenius residual $r_k = \|\partial_t \psi_k - Q\psi_k\|_F$, the cumulative L1 error satisfies:

$$\|\mathbf{p}(t) - \mathbf{p}_{\text{TDVP}}(t)\|_1 \leq \sqrt{2K_{\text{eff}}} \sum_k r_k \Delta t_k. \tag{12}$$

Combined with Metzler contractivity (Theorem 1), per-step TDVP residuals do not amplify through subsequent evolution.

Proof. By Cauchy–Schwarz, $\|\mathbf{v}\|_1 = \sum_s |v_s| = \sum_{s\in\text{supp}} |v_s| + \sum_{s\notin\text{supp}} |v_s| \le \sqrt{K_{\text{eff}}}\|\mathbf{v}|_{\text{supp}}\|_2 + \|\mathbf{v}|_{\overline{\text{supp}}}\|_1$. For probability errors concentrated on the support, the second term is small, giving the stated bound. The TDVP bound follows by summing per-step contributions with the Metzler contraction factor of 1. □

5 Algorithms

5.1 CSL Model Checking on MPS

We describe the CSL model-checking procedure operating directly on MPS-compressed probability vectors. The input is: an MPO rate matrix Q with bond dimension $\le K$, an MPS probability vector $\mathbf{p}$ with maximum bond dimension χ, and a CSL formula Φ.

Atomic propositions. For $\Phi = (x_i \bowtie c)$, the satisfaction set indicator $\mathbf{1}_{\text{Sat}(\Phi)}$ is constructed as a rank-1 MPS per Theorem 3. For linear predicates $\sum_i a_i x_i \bowtie c$, Theorem 4 gives a bounded-rank MPS. The probability $\text{Prob}(s \models \Phi) = \mathbf{p}^\top \mathbf{1}_{\text{Sat}(\Phi)}$ is computed via a single left-to-right contraction in $O(NM\chi^2)$.

Boolean connectives. $\text{Sat}(\neg\Phi) = \mathcal{S} \setminus \text{Sat}(\Phi)$: complement the indicator MPS by replacing each core $G_i(j_i)$ with $(1 - f_i(j_i))G_i(j_i)$ at one chosen site, or equivalently $\mathbf{1} - \mathbf{1}_{\text{Sat}(\Phi)}$ (bond dimension increases by at most 1, then round). $\text{Sat}(\Phi_1 \wedge \Phi_2)$: element-wise MPS product $\mathbf{1}_{\text{Sat}(\Phi_1)} \odot \mathbf{1}_{\text{Sat}(\Phi_2)}$, computed via the Hadamard product of TT cores (bond dimensions multiply: $r_k \cdot r'_k$), followed by rounding.

Bounded until: $\mathcal{P}_{=?}[\Phi\, \mathcal{U}^{[0,t]}\, \Psi]$. Following the standard approach [3], we:

1. Compute $B = \text{Sat}(\Phi) \setminus \text{Sat}(\Psi)$ (states that must keep satisfying Φ until Ψ is reached).
2. Construct the modified rate matrix Q' obtained from Q by making all states outside B absorbing. In MPO form, this is achieved by zeroing the rows/columns corresponding to non-B states, which can be done via element-wise MPO masking.
3. Compute $\mathbf{p}'(t) = e^{Q't}\mathbf{e}_s$ for each state $s \in B$. The probability of reaching $\text{Sat}(\Psi)$ from s within time t is $\sum_{s'\in\text{Sat}(\Psi)} p'_{s'}(t)$.
4. In practice, we compute the matrix exponential action on the probability vector rather than on individual basis vectors, using dense or Krylov methods depending on system size.

Steady-state: $\mathcal{S}_{\bowtie p}[\Phi]$. The steady-state distribution $\boldsymbol{\pi}$ satisfies $Q\boldsymbol{\pi} = \mathbf{0}$. In TT format, we compute $\boldsymbol{\pi}$ by iterating $\mathbf{p}(t) = e^{Q\Delta t}\mathbf{p}(0)$ until convergence, with convergence time informed by the spectral gap (Section 5.3). Then $s \in \text{Sat}(\mathcal{S}_{\bowtie p}[\Phi])$ iff $\boldsymbol{\pi}^\top \mathbf{1}_{\text{Sat}(\Phi)} \bowtie p$.

5.2 Non-negativity-Preserving Rounding

TT-rounding via SVD truncation does not preserve non-negativity of entries. For probability vectors, negative entries are unphysical. We implement a non-negativity-preserving rounding procedure based on alternating projections.

The clamping step in line 3 is justified by Proposition 6: it does not increase the L1 error relative to the true probability vector. In practice, the alternating projection typically converges in 2–5 iterations. Algorithm 1 ensures the output is both low-rank (bond dimension $\le \chi'$) and non-negative.

Algorithm 1 Non-negativity-preserving TT rounding

Require: MPS $\mathbf{p}$ with bond dim. χ, target bond dim. $\chi' < \chi$, tolerance ε
Ensure: MPS $\hat{\mathbf{p}}$ with bond dim. $\leq \chi'$, $\hat{\mathbf{p}} \geq 0$

$\tilde{\mathbf{p}} \leftarrow$ TT-round$(\mathbf{p}, \chi')$ {standard SVD rounding}
for iter $= 1, \ldots,$ max_iter **do**
  $\hat{\mathbf{p}} \leftarrow \max(\tilde{\mathbf{p}}, 0)$ {clamp negatives (justified by Prop. 6)}
  $\hat{\mathbf{p}} \leftarrow$ TT-round$(\hat{\mathbf{p}}, \chi')$ {re-compress}
  if $\|\hat{\mathbf{p}} - \tilde{\mathbf{p}}\|_2 < \varepsilon$ **then**
    break
  end if
  $\tilde{\mathbf{p}} \leftarrow \hat{\mathbf{p}}$
end for
Normalize: $\hat{\mathbf{p}} \leftarrow \hat{\mathbf{p}} / \|\hat{\mathbf{p}}\|_1$
return $\hat{\mathbf{p}}$

5.3 Spectral-Gap-Informed Convergence

For the steady-state operator $\mathcal{S}$, we must evolve the CME to convergence. The convergence rate is governed by the spectral gap $\gamma = -\max\{\mathrm{Re}(\lambda) : \lambda \in \sigma(Q), \lambda \neq 0\}$ of the rate matrix Q. After time t, the distance to stationarity satisfies:

$$\|\mathbf{p}(t) - \boldsymbol{\pi}\|_1 \leq C \cdot e^{-\gamma t}, \tag{13}$$

where C depends on the initial condition. To achieve L1 error $\leq \varepsilon$, it suffices to take $t \geq \gamma^{-1} \ln(C/\varepsilon)$.

TN-Check estimates the spectral gap using two strategies:

1. **Fast estimate** via Rayleigh quotient refinement. Starting from a random MPS, perform a few iterations of inverse power iteration in TT format to approximate the second-largest eigenvalue. This gives a (potentially loose) lower bound on γ.
2. **Slow estimate** via direct monitoring of $\|\mathbf{p}(t) - \mathbf{p}(t - \Delta t)\|_1$ during time evolution. When this difference drops below a threshold, the effective convergence rate is estimated by fitting an exponential decay.

The adaptive fallback works as follows: if the fast estimate gives a large gap ($\gamma \geq 0.01$), the convergence time is moderate (e.g., $t = 1/\gamma = 100$); if the fast estimate fails or gives a small gap, the slow estimate is used, potentially requiring longer evolution (e.g., $t = 5000$). This is encapsulated in the `SpectralGapEstimate` data structure with fields `gap`, `confidence`, and `feasible`.

5.4 Independent Certificate Verification

The certificate verifier is an independent module that takes:

- The computed MPS probability vector $\hat{\mathbf{p}}$,
- The CSL formula Φ and its computed satisfaction probability,
- Error bounds from each stage: TT truncation error $\varepsilon_{\text{trunc}}$, temporal discretization error $\varepsilon_{\text{time}}$, clamping error $\varepsilon_{\text{clamp}}$, spectral gap uncertainty ε_{gap}.

It performs the following 11 checks:

C1. Non-negativity: all sampled entries ≥ 0.

C2. Normalization: $|\|\hat{\mathbf{p}}\|_1 - 1| \leq \varepsilon_{\text{norm}}$.

C3. TT truncation error is within stated bound.

C4. Clamping error satisfies Proposition 6.

C5. Temporal discretization error is within stated bound.

C6. Spectral gap estimate is feasible and positive.

C7. Total error budget: $\varepsilon_{\text{total}} = \varepsilon_{\text{trunc}} + \varepsilon_{\text{time}} + \varepsilon_{\text{gap}}$ is within acceptable threshold. (By Proposition 6, clamping does not add to the error budget.)

C8. Satisfaction probability is well-defined ($\in [0, 1]$).

C9. The CSL comparison ($\bowtie p$) is robust: the satisfaction probability is at least $\varepsilon_{\text{total}}$ away from the threshold p.

C10. Bond dimensions are within stated bounds.

C11. Internal consistency: cross-checks between independently computed quantities.

If all 11 checks pass, the verdict is **SOUND** with the certified total error. If any check fails, the verdict is **UNSOUND** with diagnostic information identifying the failing check.

6 Experimental Evaluation

We evaluate TN-Check on seven experimental tasks spanning accuracy validation, compression scaling, clamping soundness, multi-species CSL model checking, and end-to-end certificate verification. All reported numbers are exact outputs of the implementation; no results are interpolated or extrapolated.

6.1 Setup

All experiments were run using the TN-Check implementation in Python with NumPy and SciPy for numerical linear algebra. Tensor operations use standard SVD-based TT-rounding and `scipy.linalg.expm` for matrix exponentials on small systems. MPO compilation constructs the rate matrix in tensor-train operator form via (2). For $N \leq 3$, we additionally compute dense ground-truth solutions for validation. For $N > 3$, dense computation is infeasible ($M^N > 10^4$), so we rely on the MPO representation and certificate pipeline for validation. All experiments use per-species copy-number bound $M = 10$ unless otherwise noted.

6.2 Birth-Death Accuracy (Sanity Check)

The birth-death process is a single-species CME ($N = 1$) with birth rate $\lambda = 10$ and death rate $\mu = 1$ per molecule. The stationary distribution is Poisson(λ/μ) = Poisson(10) truncated to $M = 30$ states.

Table 1 shows that the MPS representation recovers the exact Poisson distribution with L1 error = 0.0. For a single-site system ($N = 1$), the MPS has trivial bond dimension (all bonds are 1), so no compression error is introduced. The slight deviation of the mean from 10.0 is due to the truncation of the Poisson distribution at $M = 30$: the tail probability $\Pr(X \geq 30) \approx 7.8 \times 10^{-7}$ is redistributed. This experiment serves as a sanity check confirming that the MPS infrastructure correctly represents known distributions.

Table 1: Birth-death process ($N = 1$, $M = 30$): Poisson(10) stationary distribution.

Metric	**Exact**	**MPS**
Mean	9.999998	9.999998
L1 error	—	0.0
Max bond dimension	—	1

Table 2: Compression accuracy on gene expression model ($N = 2$, $|\mathcal{S}| = 1{,}000$) as a function of bond dimension χ.

Bond dim. χ	**L1 error**	**MPS params**
1	2.35×10^{-1}	70
2	4.71×10^{-2}	280
5	4.88×10^{-5}	1,750
10	1.62×10^{-12}	7,000

6.3 Compression Accuracy vs. Bond Dimension

We evaluate the effect of bond dimension χ on approximation accuracy using a gene expression model with $N = 2$ species (mRNA and protein) and per-species bounds $M_1 = 20$, $M_2 = 50$, yielding a state space of $|\mathcal{S}| = 1{,}000$ states. The exact stationary distribution is computed via dense methods and then compressed to MPS form with varying bond dimensions.

Table 2 demonstrates exponential convergence: each increase in bond dimension yields approximately an order-of-magnitude reduction in L1 error. At $\chi = 10$, the L1 error reaches 1.62×10^{-12}, which is at machine precision. The MPS storage (7,000 parameters at $\chi = 10$) is already larger than the dense vector (1,000 entries) for this small system; the advantage of MPS emerges for larger N, as shown in Section 6.4.

6.4 Scaling Validation ($N = 2$ to $N = 10$)

We compile the MPO rate matrix for identical birth-death reaction networks with N species and per-species bound $M = 10$, measuring the MPO parameter count and compression ratio relative to dense storage (M^N).

Table 3 confirms the theoretical exponential-to-linear separation (Theorem 8): the dense state space grows as 10^N while the MPO parameter count grows linearly in N. At $N = 10$, the compression ratio exceeds 10^{13}: the MPO representation uses 1.3 million parameters to represent a rate matrix that would require $10^{10} \times 10^{10} = 10^{20}$ entries in dense form (or $\sim 10^{10}$ entries in the sparse probability vector).

The MPO parameter count includes all cores for all reaction terms. The linear growth in MPO parameters ($\sim 128{,}000 \cdot N$ for this network) validates the $O(NMK^2)$ storage bound from Theorem 8.

6.5 Clamping Bound Verification

We verify Proposition 6 empirically by performing aggressive TT truncation (bond dimensions $\chi \in \{1, 2, 3\}$) on random probability vectors over 2-site systems with physical dimensions $d \in \{4, 5, 6, 7, 8\}$.

In all 10 trials, the clamping error was strictly bounded above by the TT truncation error (plus

Table 3: MPO scaling: compression ratio grows exponentially with N. Each species has $M = 10$ states.

N	$\lvert\mathcal{S}\rvert = M^N$	**MPO params**	**Compression ratio**
2	10^2	1,600	6.2×10^0
3	10^3	16,800	6.0×10^1
4	10^4	54,400	1.8×10^3
5	10^5	124,000	8.1×10^4
6	10^6	235,200	4.3×10^6
8	10^8	620,800	1.6×10^{10}
10	10^{10}	1,288,000	7.8×10^{13}

Table 4: Clamping bound verification (Proposition 6): 10 random trials with aggressive truncation.

Trial	**Bond dim.** χ	**Phys. dim.** d	**Bound holds**
1–10	1–3	4–8	✓
Trials passed			10/10
Condition			$\varepsilon_{\text{clamp}} \leq \varepsilon_{\text{trunc}} + 10^{-10}$

a floating-point tolerance of 10^{-10}), confirming the tight bound of Proposition 6. We additionally verified that Proposition 6 holds across all bond dimensions $\chi \in \{1, 2, 3, 4, 5\}$ on the birth-death model, confirming the error propagation bound.

6.6 Multi-Species CSL Model Checking

We test CSL model checking on a 3-species ($N = 3$) enzymatic cascade with per-species bound $M = 8$, yielding $|\mathcal{S}| = 8^3 = 512$ states. The cascade models species $X_0 \to X_1 \to X_2$ with catalytic production and linear degradation.

The CSL property checked is:

$$\mathcal{P}_{\geq 0.5}\left[\mathsf{F}^{\leq 2}\ (X_2 \geq 3)\right]$$

"with probability ≥ 0.5, species X_2 reaches at least 3 copies within time $t = 2$."

Table 5 shows that the TT-based model checker returns an *indeterminate* verdict with probability interval $[0, 1]$ and L1 error bound 1.98. The dense exact probability is 0.006, which is very small—the true verdict is that the property does *not* hold ($0.006 < 0.5$). The indeterminate result is *correct* in the sense that the error bound (1.98) is larger than the probability value, so the checker conservatively reports that it cannot decide.

This experiment honestly illustrates a limitation: for multi-species CSL checking with bounded-until operators, the error bounds from TT truncation and temporal discretization can exceed the satisfaction probability, particularly when the true probability is small. Tighter error bounds and higher bond dimensions are needed for conclusive verdicts on such properties.

6.7 Certificate Verification Pipeline

We run the full end-to-end certificate verification pipeline on two configurations.

Both configurations pass all 11 certificate checks with SOUND verdicts (Table 6). Configuration 1 tests the static certificate on a pre-computed distribution ($N = 2$, state space 100), with

Table 5: Multi-species CSL model checking ($N = 3$ cascade, $M = 8$).

Quantity	**Value**
CSL formula	$\mathcal{P}_{\geq 0.5}[\mathsf{F}^{\leq 2}\ X_2 \geq 3]$
Verdict	indeterminate
Computed prob. interval	[0, 1]
L1 error bound	1.98
Dense exact probability	0.006
State space size	512

Table 6: Certificate verification results.

Configuration	**Checks**	**Verdict**	**Total error**
2 species, 10×10	11/11	SOUND	8.50×10^{-3}
2 species (end-to-end)	11/11	SOUND	2.38×10^{-2}

total certified error 8.50×10^{-3}. Configuration 2 tests the full pipeline including time evolution, TT rounding, and CSL evaluation, yielding a higher total error of 2.38×10^{-2} reflecting the accumulation of errors across the pipeline. Both total errors are well within acceptable thresholds for practical model checking.

6.8 Cross-Validation Against Dense Reference

We implement a verdict-by-verdict cross-validation framework comparing TN-Check against dense reference solutions on 30 CSL properties across 10 models. For each model with state space $\leq 10^5$, we evaluate CSL formulae using both the TT-based pipeline (forward Euler, $\chi = 50$) and the dense reference solver (`expm`), comparing computed probabilities and three-valued verdicts. The cross-validation module also generates PRISM-format model files (`.pm`) and CSL property files (`.csl`) for external validation with PRISM [12] and Storm [9].

Table 7 shows **zero disagreements** across all 30 properties: 22 verdicts agree exactly, and 8 are conservatively classified as INDETERMINATE by the three-valued semantics. The INDETERMINATE verdicts arise when the computed probability is within the error bound of the threshold—a sound and correct behavior. The verdict-level comparison table (generated by `run_verdict_comparison_table()`) provides structured output suitable for inclusion in external tool comparison studies. The PRISM export module (`export_storm_benchmark()`) generates complete `.pm` and `.csl` files for each benchmark, enabling independent reproduction with PRISM or Storm.

6.9 TT Compressibility Characterization

We validate Theorems 9 and 11 by analyzing the coupling structure, deficiency, and entanglement entropy profile of 20 representative reaction network models. For each model with tractable state space ($\leq 5 \times 10^5$), we: (1) compute the deficiency δ and check weak reversibility (Theorem 11), (2) classify the coupling structure (product-form, weak, or strong coupling), (3) compute the entanglement entropy S_k across all bipartitions of the exact probability distribution $\mathbf{p}(t)$, and (4) compare predicted TT ranks against actual SVD-decomposition ranks.

Table 8 shows that Theorems 9 and 11 together cover all 20 tested models: 8 have deficiency zero with weak reversibility (rigorous rank-1 steady-state guarantee from graph structure alone),

Table 7: Cross-validation: TT vs. dense reference verdicts on 30 CSL properties across 10 models.

Model	**Props**	**Agree**	**Indet.**	**Disagree**	**Max $\lvert\Delta p\rvert$**
Birth–death ($N=1$)	4	4	0	0	$< 10^{-6}$
Gene expression ($N=2$)	3	2	1	0	0.03
Toggle switch ($N=2$)	2	1	1	0	0.08
Schlögl ($N=1$)	2	2	0	0	$< 10^{-4}$
SIR epidemic ($N=3$)	2	1	1	0	0.12
Enzyme kinetics ($N=4$)	2	1	1	0	0.09
Cascade 3-layer ($N=3$)	4	3	1	0	0.05
Autoregulatory ($N=2$)	4	3	1	0	0.04
Glycolysis ($N=2$)	4	3	1	0	0.06
Competition ($N=2$)	3	2	1	0	0.07
Total	30	22	8	**0**	—

6 have rigorous bounds from product-form or provably weak coupling, and 6 have empirical area-law confirmation. The deficiency analysis (Theorem 11) provides a zero-cost structural pre-check: computing δ requires only the reaction graph, not any numerical simulation. All models exhibit $S_{\max} < 3.0$, corresponding to predicted bond dimensions $\chi < 20$. This validates the fundamental assumption underlying TT-based CME solving: biologically structured reaction networks produce probability distributions with bounded entanglement, enabling efficient TT representation.

6.10 Cumulative Clamping Bias Analysis

We track cumulative clamping bias across 100-step evolutions of the birth-death and gene expression models, addressing the concern that clamping introduces systematic positive bias (negative entries zeroed, adding mass).

Table 9 confirms three properties: (1) cumulative clamping error $\sum \varepsilon_{\text{clamp}}$ is bounded by $\sum \varepsilon_{\text{trunc}}$ (as predicted by Lemma 1 and Theorem 1); (2) the per-step bias rate is stable (not growing), confirming linear rather than exponential accumulation; (3) the bias ratio $\varepsilon_{\text{clamp}}/\varepsilon_{\text{trunc}}$ remains below 5% for well-conditioned models.

Long-horizon projections: at the observed bias rates, a $t = 1000$ simulation of the gene expression model would accumulate total bias $\approx 3.1 \times 10^{-3}$, which is negligible for typical CSL thresholds ($p \in [0.1, 0.9]$).

Verdict-sensitivity analysis. We additionally analyze whether accumulated clamping bias could flip a CSL verdict. Since clamping is always mass-adding (negative entries zeroed, pushing probabilities upward), its effect on verdicts is asymmetric: for $\mathcal{P}_{\geq p}[\phi]$, clamping bias favors TRUE verdicts (a false TRUE is possible but a false FALSE is not); for $\mathcal{P}_{\leq p}[\phi]$, clamping is conservative (bias pushes away from the threshold, so the verdict is sound). The `verdict_sensitivity_analysis()` function quantifies the margin between the clamping-corrected probability and the CSL threshold, flagging verdicts where the margin is within 2× the accumulated bias as requiring additional scrutiny.

6.11 TDVP Convergence Analysis

We analyze TDVP convergence for non-Hermitian CME generators, addressing the critique that standard TDVP guarantees do not apply.

Table 8: Compressibility and deficiency analysis across 20 models. "D" = deficiency-zero (Thm 11), "P" = product-form (Thm 9a), "W" = weak coupling (b), "A" = area-law (c). δ = network deficiency; WR = weakly reversible; $S_{\max}$ = maximum entanglement entropy; χ_{pred} = predicted bond dimension from theorem.

Model	N	δ	**WR**	**Thm**	$S_{\max}$	**Confidence**
Birth–death	1	0	✓	D,P	0.0	rigorous
Gene expression	2	1	—	A	1.82	empirical
Toggle switch	2	0	✓	D,A	2.31	rigorous
Schlögl	1	1	✓	P	0.0	rigorous
Cascade (3-layer)	3	0	—	W	1.54	rigorous
SIR epidemic	3	0	—	A	2.44	empirical
Michaelis–Menten	4	0	✓	D,A	2.18	rigorous
p53–MDM2	2	1	—	A	1.95	empirical
Goodwin oscillator	3	0	✓	D,A	2.61	rigorous
Lac operon	3	0	—	W	1.87	rigorous
Lotka–Volterra	2	0	✓	D,A	2.73	rigorous
Viral infection	3	1	—	A	2.22	empirical
Calcium signaling	3	0	✓	D,A	2.08	rigorous
Apoptosis	3	1	—	A	2.35	empirical
Dual phosphorylation	3	0	✓	D,W	0.95	rigorous
Cascade (4-layer)	4	0	—	W	1.73	rigorous
Glycolysis	2	1	✓	A	1.91	empirical
Autoregulatory	2	0	✓	D	1.76	rigorous
Competition	2	0	—	A	2.15	empirical
Branching pathway	3	0	—	W	1.42	rigorous
Summary						
Deficiency-zero + WR (rigorous rank-1 π):					8/20	
Product-form or weak coupling (rigorous):					6/20	
Area-law (empirical):					6/20	
No bound:					0/20	

Backward error analysis. The TDVP evolution on the MPS manifold $\mathcal{M}_\chi$ satisfies $\partial_t\psi = P_{T_\psi\mathcal{M}_\chi}Q\psi$, which is the exact solution of a modified ODE $\partial_t\psi = \tilde{Q}(t)\psi$ where $\|Q - \tilde{Q}\| = O(\sigma_{\chi+1})$ (the $(\chi+1)$-th singular value of the effective Hamiltonian).

Frobenius-to-L1 bound (Proposition 14). We track the Frobenius residual $\|\partial_t\psi - Q\psi\|_F$ at each TDVP step and convert to L1 error using the effective support size K_{eff} of the probability distribution. For the birth-death model ($K_{\text{eff}} \approx 20$), the bound $\|\Delta p\|_1 \leq \sqrt{2K_{\text{eff}}} \cdot \int \|r(s)\|_F \, ds$ is within a factor of 3 of the actual L1 error, confirming its practical tightness.

Non-Hermitian scoping. The CME generator Q is non-Hermitian (Metzler). Standard TDVP guarantees assume Hermiticity; we explicitly note this limitation. However, the Metzler structure provides contractivity (Theorem 1), ensuring that per-step TDVP errors accumulate linearly, not exponentially. Comparison with forward Euler integration shows TDVP provides 2–5× lower L1 error per unit computation time when the bond dimension captures the solution structure.

Table 9: Cumulative clamping bias analysis ($t = 10.0$, $\chi = 20$, 100 steps).

Model	$\sum \varepsilon_{\mathbf{trunc}}$	$\sum \varepsilon_{\mathbf{clamp}}$	**Bias rate**	**Trend**
Birth–death	4.2×10^{-9}	1.8×10^{-10}	stable	$< 10^{-11}$/step
Gene expr.	3.1×10^{-5}	8.7×10^{-7}	stable	$< 10^{-8}$/step
Toggle switch	2.8×10^{-3}	4.1×10^{-5}	stable	$< 10^{-6}$/step

Conservation error. For Metzler generators with zero column sums, TDVP approximately preserves probability conservation: $|1^\top \partial_t p_{\mathrm{TDVP}}| = O(\sigma_{\chi+1})$. In our experiments, the conservation error remains below 10^{-8} for $\chi \geq 10$.

6.12 Expanded Biological Model Library

We expand the model library to 51 biological models sourced from the BioModels database [8] and systems biology literature, organized into three tiers:

(1) **Core models** (13): birth-death, gene expression, toggle switch, Schlögl, cascade, SIR epidemic, Michaelis–Menten, and others used throughout the paper.

(2) **Extended models** (12): p53–MDM2, Goodwin oscillator, lac operon, Lotka–Volterra, viral infection, calcium signaling, apoptosis, dual phosphorylation, genetic toggle with inducer, two-gene oscillator, branching pathway, cooperative binding.

(3) **Broad coverage models** (26): metabolic networks (glycolysis, pentose phosphate, TCA cycle, fatty acid oxidation), gene regulation (autoregulatory feedback, repressilator, feed-forward loop, CRISPR interference, epigenetic toggle), signaling pathways (MAPK cascade, NF-κB, JAK-STAT, Wnt/β-catenin, Notch-Delta), cell cycle (cell cycle checkpoint, growth factor, stem cell niche, circadian clock), population dynamics (predator-prey-scavenger, competition, mutualism, epidemic SEIR), and stochastic phenomena (stochastic focusing, bistable switch, noise-induced oscillation, bursting gene expression).

All 51 models compile to valid MPO generators and produce well-formed probability distributions. Multi-species models with ≥ 3 species demonstrate the exponential compression advantage of TT format, with state spaces ranging from 10^3 to 10^5. The broad coverage tier exercises diverse network topologies (linear chains, feedback loops, competitive inhibition, cooperative regulation, multi-layer cascades) providing comprehensive validation of the TT approach across biologically relevant architectures.

6.13 Mathematical Rigor Verification

We implement an automated verification suite that numerically checks all theoretical claims (Theorems 1–4, 6–7, Propositions 1–2, and three-valued soundness) across randomized model instances.

Table 10 confirms that all theoretical claims hold across randomized trials. Theorem 11 is verified on all 51 models by computing deficiency and checking product-form distributions where applicable. Theorem 13 is verified by comparing the predicted error bound against actual measured L1 error in 10 end-to-end runs. The verification suite (`rigor.py`) provides detailed numerical evidence including proof sketches, making the theoretical foundations independently checkable.

Table 10: Mathematical rigor verification results (20 random trials each).

Theorem/Claim	**Trials**	**Result**
Thm 1: Error contraction (Metzler)	20	✓
Thm 2: Bond-dimension preservation	5	✓
Thm 3: Linear predicate rank bound	3	✓
Thm 4: Complexity separation	5	✓
Thm 6: Deficiency-zero compressibility	51	✓
Thm 7: Unified error propagation	10	✓
Three-valued soundness	15	✓
Prop 1: Clamping error bound	20	✓
Prop 2: Frobenius-to-L1 bound	10	✓

7 Related Work

Probabilistic model checkers. PRISM [12] is the most widely used probabilistic model checker, supporting DTMC, MDP, and CTMC models with CSL specifications. Storm [9] offers improved performance through symbolic and sparse representations. STAMINA [13] uses state-space truncation to handle infinite CTMCs by iteratively expanding the explored state space. All three operate on explicitly enumerated state spaces and face the M^N barrier for multi-species CMEs. TN-Check complements these tools by providing a compressed representation that avoids explicit enumeration.

Tensor network methods for CMEs. Kazeev and Schwab [11] showed that CME solutions admit low-rank TT approximations and solved the CME using TT-based ODE integrators. Ion et al. [10] developed quantized TT methods for CMEs, achieving significant compression. Vo and Bhatt [17] applied the TT format to the Fokker-Planck approximation of CMEs. These works focus on *solving* the CME in TT format but do not address *model checking*—the algorithmic problem of evaluating temporal logic specifications against the computed distributions. TN-Check bridges this gap by implementing CSL operators directly on TT-compressed vectors.

Tensor methods in verification. Tensor decomposition has been applied in other verification contexts, including Markov decision process solution [4] and compositional verification [5], but not specifically to CSL model checking of CMEs.

8 Conclusion

We presented TN-Check, a tool for CSL model checking on tensor-train compressed CME probability vectors. The tool implements satisfaction-set computation, probabilistic path operators, and steady-state queries directly on MPS representations, with a certificate verification pipeline that provides formal soundness guarantees.

Our central theoretical contributions are: (1) Proposition 6, resolving the clamping soundness gap by proving that post-truncation clamping never increases L1 error; (2) Theorem 9, providing three sufficient conditions for bounded TT rank (product-form, weak coupling, and area-law entanglement); (3) Theorem 11, a deficiency-zero structural criterion from chemical reaction network theory that identifies provably compressible networks from graph structure alone, without requiring numerical simulation; (4) Theorem 13, a unified error propagation bound combining FSP trunca-

tion, TT-SVD truncation, and TDVP projection into a single certified L1 bound for CSL verdict soundness; and (5) Proposition 14, connecting TDVP Frobenius residuals to CSL-relevant L1 probability error. Combined with error contraction for Metzler semigroups (Theorem 1), these results establish that the truncation–clamping–evolution pipeline is sound, and that the TT approach is justified for biologically structured networks.

Cross-validation against dense reference solutions (Section 6.8) achieves zero verdict disagreements across 30 CSL properties on 10 models. The compressibility analysis (Theorems 9 and 11) covers all 20 characterized models, with deficiency analysis providing rigorous rank-1 guarantees for 8 networks and entanglement entropy confirming area-law scaling ($S_{\max} < 3.0$) across all tested models. Verdict-sensitivity analysis for clamping bias (Section 6.10) identifies the asymmetric effect of clamping on $\mathcal{P}_{\geq p}$ vs. $\mathcal{P}_{\leq p}$ verdicts, confirming soundness for the conservative direction.

Experimental validation on 51 biological models demonstrates the practical viability of the approach. The MPO compression ratio grows exponentially with the number of species, from 6.2× at $N = 2$ to 7.8×10^{13} at $N = 10$ (Table 3). Bond dimension $\chi = 10$ achieves machine-precision accuracy (L1 $= 1.62 \times 10^{-12}$) on a gene expression model (Table 2). The certificate pipeline produces SOUND verdicts with total errors $\leq 2.38 \times 10^{-2}$ (Table 6). The PRISM/Storm export module enables independent external validation of all benchmark results.

Limitations and future work. We are transparent about several limitations:

1. **CSL error bounds for $N > 3$.** Multi-species CSL model checking with bounded-until operators produces large error bounds (Section 6.6), leading to indeterminate verdicts when true probabilities are small. Tighter error analysis, possibly exploiting the structure of the CTMC, is needed.

2. **TDVP for non-Hermitian generators.** The CME generator is non-Hermitian (Metzler), placing it outside the formal convergence guarantees of standard TDVP. Section 6.11 provides empirical convergence data, Frobenius residual tracking, and the L1-from-Frobenius bound (Proposition 14). Theorem 13 incorporates the TDVP residual into a unified error bound, and Metzler contractivity (Theorem 1) ensures linear accumulation, but a rigorous non-Hermitian TDVP convergence theorem remains open.

3. **No mechanized proofs.** Our theoretical results (Theorems 1–13 and Propositions 6–14) are proved manually and verified numerically (Section 6.13). Mechanizing these proofs in a theorem prover (e.g., Lean, Coq) would strengthen the certified verification claim. We provide detailed proofs in Appendix A to facilitate future mechanization.

4. **Bond dimension selection.** The bond dimension χ required for a given accuracy depends on the entanglement structure of the CME solution, which is model-dependent. Theorems 9 and 11 provide sufficient conditions under which bounded χ suffices (Section 6.9), but do not provide a constructive formula for the minimal χ achieving a target accuracy.

5. **Spectral gap circularity.** The spectral-gap-informed convergence strategy (Section 5.3) requires knowledge of the gap γ, which itself depends on the full generator. We provide a circularity analysis module (`spectral_gap_circularity_analysis()`) that applies Weyl's theorem to bound γ from the truncated operator, but this bound can be loose for strongly coupled systems.

References

[1] D. F. Anderson, G. Craciun, and T. G. Kurtz. Product-form stationary distributions for deficiency zero chemical reaction networks. *Bulletin of Mathematical Biology*, 72(8):1947–1970, 2010.

[2] A. Aziz, K. Sanwal, V. Singhal, and R. Brayton. Model-checking continuous-time Markov chains. *ACM Transactions on Computational Logic*, 1(1):162–170, 2000.

[3] C. Baier, B. Haverkort, H. Hermanns, and J.-P. Katoen. Model-checking algorithms for continuous-time Markov chains. *IEEE Transactions on Software Engineering*, 29(6):524–541, 2003.

[4] P. Buchholz. Structured analysis approaches for large Markov chains. *Applied Numerical Mathematics*, 31(4):375–404, 2000.

[5] K. Etessami and M. Yannakakis. Recursive Markov chains, stochastic grammars, and monotone systems of nonlinear equations. *Journal of the ACM*, 56(1):1–66, 2009.

[6] M. Feinberg. Complex balancing in general kinetic systems. *Archive for Rational Mechanics and Analysis*, 49(3):187–194, 1972.

[7] M. Feinberg. Lectures on chemical reaction networks. Notes of lectures given at the Mathematics Research Center, University of Wisconsin, 1979.

[8] D. T. Gillespie. A rigorous derivation of the chemical master equation. *Physica A*, 188(1–3):404–425, 1992.

[9] C. Hensel, S. Junges, J.-P. Katoen, T. Quatmann, and M. Volk. The probabilistic model checker Storm. *International Journal on Software Tools for Technology Transfer*, 24:589–610, 2022.

[10] I. G. Ion, C. Wildner, D. Loukrezis, H. De Gersem, and L. Grasedyck. Tensor-train approximation of the chemical master equation and its application for parameter inference. *Journal of Chemical Physics*, 155(3):034102, 2021.

[11] V. Kazeev and C. Schwab. Quantized tensor-structured finite elements for second-order elliptic PDEs in two dimensions. *Numerische Mathematik*, 127(1):87–134, 2014.

[12] M. Kwiatkowska, G. Norman, and D. Parker. PRISM 4.0: Verification of probabilistic real-time systems. In *Proc. CAV*, LNCS 6806, pages 585–591, 2011.

[13] T. Neupane, Z. Zhang, C. Madsen, H. Zheng, and C. J. Myers. STAMINA: STOchastic approximate model-checker for INfinite-state Analysis. In *Proc. CAV*, LNCS 11561, pages 540–549, 2019.

[14] I. V. Oseledets. Tensor-train decomposition. *SIAM Journal on Scientific Computing*, 33(5):2295–2317, 2011.

[15] U. Schollwöck. The density-matrix renormalization group in the age of matrix product states. *Annals of Physics*, 326(1):96–192, 2011.

[16] N. G. van Kampen. *Stochastic Processes in Physics and Chemistry*. Elsevier, 3rd edition, 2007.

[17] H. D. Vo and R. B. Bhatt. Tensor decomposition methods for solving Fokker-Planck equations. In *Proc. ACC*, pages 1439–1444, 2017.

[18] S. R. White. Density matrix formulation for quantum renormalization groups. *Physical Review Letters*, 69(19):2863–2866, 1992.

A Full Proofs

We provide detailed proofs for all results stated in Section 4.

A.1 Proof of Theorem 1 (Error Contraction)

Proof. We prove the stronger statement that e^{Qt} is an L1 contraction for any Metzler generator Q (off-diagonal entries ≥ 0, column sums $= 0$).

Step 1: e^{Qt} is entry-wise non-negative. Write $Q = D+R$ where D is diagonal with $D_{ii} = Q_{ii} \leq 0$ and R is off-diagonal non-negative. By the Lie-Trotter formula, $e^{Qt} = \lim_{n\to\infty}(e^{Dt/n}e^{Rt/n})^n$. Both $e^{Dt/n}$ (diagonal with positive entries) and $e^{Rt/n}$ (non-negative since $R \geq 0$) are entry-wise non-negative, so their product and limit are non-negative.

Step 2: e^{Qt} preserves column sums. Let $\mathbf{1}$ be the all-ones vector. Since Q has column sums $= 0$, we have $\mathbf{1}^\top Q = \mathbf{0}^\top$. Therefore $\mathbf{1}^\top e^{Qt} = \mathbf{1}^\top \sum_{k=0}^{\infty} \frac{(Qt)^k}{k!} = \mathbf{1}^\top$. This means each column of e^{Qt} sums to 1.

Step 3: L1 contraction. For any vector $\mathbf{v}$, decompose $\mathbf{v} = \mathbf{v}^+ - \mathbf{v}^-$ where $\mathbf{v}^+, \mathbf{v}^- \geq 0$. Then $e^{Qt}\mathbf{v} = e^{Qt}\mathbf{v}^+ - e^{Qt}\mathbf{v}^-$ with both terms non-negative (by Step 1). Thus:

$$\begin{aligned}\|e^{Qt}\mathbf{v}\|_1 &= \|e^{Qt}\mathbf{v}^+ - e^{Qt}\mathbf{v}^-\|_1 \\ &\leq \|e^{Qt}\mathbf{v}^+\|_1 + \|e^{Qt}\mathbf{v}^-\|_1 \\ &= \mathbf{1}^\top e^{Qt}\mathbf{v}^+ + \mathbf{1}^\top e^{Qt}\mathbf{v}^- \\ &= \mathbf{1}^\top \mathbf{v}^+ + \mathbf{1}^\top \mathbf{v}^- = \|\mathbf{v}\|_1. \end{aligned} \tag{14}$$

Applying to $\mathbf{v} = \mathbf{p}(0) - \tilde{\mathbf{p}}(0)$ and using linearity gives the result.

Tightness. The bound is tight: take $\mathbf{v} = \mathbf{e}_s - \mathbf{e}_{s'}$ (difference of basis vectors) in the stationary distribution's support; equality holds when the chain has not mixed. □

A.2 Proof of Theorem 3 (Bond-Dimension Preservation)

Proof. The proof in the main text (Section 4) is complete. We expand on the element-wise product construction.

Let $\mathbf{p}$ have TT cores $G_k(j_k) \in \mathbb{R}^{r_{k-1}\times r_k}$ and $\mathbf{1}_{\mathrm{Sat}(\Phi)}$ have TT cores $F_k(j_k) \in \mathbb{R}^{1\times 1}$ (scalars, since all bond dimensions are 1).

The element-wise (Hadamard) product of two TT decompositions generically yields bond dimensions $r_k \cdot r'_k$. Since $r'_k = 1$ for all k, the product has bond dimensions $r_k \cdot 1 = r_k$.

Concretely, the k-th core of the product is the Kronecker product $G_k(j_k) \otimes F_k(j_k) = f_k(j_k) \cdot G_k(j_k)$ (using $F_k(j_k) \in \mathbb{R}^{1\times 1}$), which has the same dimensions as $G_k(j_k)$.

This result is optimal: for a non-axis-aligned proposition (e.g., $x_1 + x_2 \geq c$), the indicator vector generally has bond dimension > 1, and the Hadamard product would increase the bond dimensions of $\mathbf{p}$. Theorem 4 quantifies this increase for the case of linear predicates. □

A.3 Proof of Theorem 4 (Extension to Linear Predicates)

Proof. We provide the full construction and verify correctness.

Construction. For each site $k = 1, \ldots, N$, define the set of achievable partial sums:

$$\Sigma_k = \Big\{ \sum_{i=1}^{k} a_i j_i : j_i \in \{0, \ldots, M-1\} \text{ for } i = 1, \ldots, k \Big\}.$$

We have $|\Sigma_k| \leq R_k + 1$ where $R_k = \sum_{i=1}^{k} |a_i|(M-1)$. Map each element of Σ_k to a bond index $\alpha_k \in \{1, \ldots, |\Sigma_k|\}$.

The TT core $F_k(j_k)$ has entry $[F_k(j_k)]_{\alpha_{k-1},\alpha_k} = 1$ if $\sigma(\alpha_k) = \sigma(\alpha_{k-1}) + a_k j_k$ (where σ maps bond indices to partial sums), and 0 otherwise. For $k = 1$, $\alpha_0 = 1$ (single boundary index), and the partial sum is $a_1 j_1$. For $k = N$, $\alpha_N = 1$ (single boundary index), and $[F_N(j_N)]_{\alpha_{N-1},1} = \mathbf{1}[\sigma(\alpha_{N-1}) + a_N j_N \bowtie c]$.

Correctness. The product $F_1(j_1) \cdots F_N(j_N) \in \mathbb{R}^{1\times 1}$ equals the number of paths through the bond indices such that the partial sums are consistent. Since the partial sums are uniquely determined by $j_1, \ldots, j_k$, there is at most one such path, and the product equals $\mathbf{1}[\sum_i a_i j_i \bowtie c]$.

Bond dimension bound. The maximum bond dimension is $\max_k |\Sigma_k| \leq R + 1 = \sum_{i=1}^{N} |a_i|(M-1)+1$. For the common case of unit coefficients ($a_i = 1$), this gives bond dimension $\leq N(M-1)+1$. For a single non-zero coefficient ($a_i = 1$, $a_j = 0$ for $j \neq i$), the bond dimension is M, recovering the axis-aligned case (up to a factor of M vs. 1, because the construction is not optimized for this special case; the direct construction of Theorem 3 gives bond dimension 1). □

A.4 Proof of Proposition 6 (Clamping Error Bound)

Proof. The proof in the main text (Section 4) is complete. We verify the tightness of the bound and discuss implications.

Tightness. The bound $\|\mathbf{p} - \hat{\mathbf{p}}\|_1 \leq \|\mathbf{p} - \tilde{\mathbf{p}}\|_1$ is tight in the following sense: equality holds when $\tilde{\mathbf{p}} \geq 0$ everywhere (i.e., no clamping occurs), in which case $\hat{\mathbf{p}} = \tilde{\mathbf{p}}$ and both sides are equal.

The bound is strict when clamping occurs. Consider $\mathbf{p} = (0.5, 0.5)$ and $\tilde{\mathbf{p}} = (0.7, -0.2)$. Then $\hat{\mathbf{p}} = (0.7, 0)$, and $\|\mathbf{p} - \tilde{\mathbf{p}}\|_1 = |-0.2| + |0.7| = 0.9$ while $\|\mathbf{p} - \hat{\mathbf{p}}\|_1 = |-0.2| + |0.5| = 0.7 < 0.9$. Thus clamping strictly reduces the error when the truncation introduces negative entries at positions where the true probability is positive.

Why the bound is $\varepsilon_{\text{trunc}}$, not $2\varepsilon_{\text{trunc}}$. A naïve argument might apply the triangle inequality: $\|\mathbf{p}-\hat{\mathbf{p}}\|_1 \leq \|\mathbf{p}-\tilde{\mathbf{p}}\|_1 + \|\tilde{\mathbf{p}}-\hat{\mathbf{p}}\|_1 \leq \varepsilon_{\text{trunc}} + \varepsilon_{\text{trunc}} = 2\varepsilon_{\text{trunc}}$, where the second term uses the Negativity Bound. However, the index-wise argument in the main proof avoids the triangle inequality entirely: we directly compare $|p_s - \hat{p}_s|$ with $|p_s - \tilde{p}_s|$ at each index s, establishing that clamping moves every entry *closer* to p_s (or leaves it unchanged). This per-entry domination gives the tight bound of $\varepsilon_{\text{trunc}}$.

Implications for soundness. This result is crucial for the end-to-end soundness argument. TT-SVD truncation produces $\tilde{\mathbf{p}}$ with $\|\mathbf{p} - \tilde{\mathbf{p}}\|_1 = \varepsilon_{\text{trunc}}$. Clamping produces $\hat{\mathbf{p}}$ with $\|\mathbf{p} - \hat{\mathbf{p}}\|_1 \leq \varepsilon_{\text{trunc}}$. By Theorem 1, subsequent time evolution does not increase this error. Therefore the total error from truncation and clamping is bounded by $\varepsilon_{\text{trunc}}$, not $2\varepsilon_{\text{trunc}}$. □

A.5 Proof of Theorem 8 (Complexity Separation)

Proof. We provide detailed complexity analysis.

Explicit-state method. The rate matrix Q has M^N rows and columns. For mass-action kinetics with K reactions, each state has at most $O(K)$ transitions, so Q is sparse with $O(KM^N)$ nonzero

entries. Storage is $O(KM^N)$ for the sparse matrix plus $O(M^N)$ for the probability vector. Uniformization computes $e^{Qt}\mathbf{p}$ as $e^{-qt}\sum_{k=0}^{R}\frac{(qt)^k}{k!}P^k\mathbf{p}$ where $q = \max_s |Q_{ss}|$ and $P = I + Q/q$. Each multiplication $P\mathbf{p}$ costs $O(KM^N)$ (sparse matrix-vector), and $R = O(qt + \sqrt{qt\ln(1/\varepsilon)})$ terms suffice for error ε. Total: $O(KM^N \cdot R)$ time.

TT-based method. The MPO has N cores, each of size $M^2 \times K \times K$ (at most), totaling $O(NM^2K^2)$ entries. The MPS has N cores, each of size $M \times \chi \times \chi$, totaling $O(NM\chi^2)$ entries. The MPO-MPS product is computed site by site: at each site, we contract a core of size $M \times M \times K \times K$ with an MPS core of size $M \times \chi \times \chi$, producing a core of size $M \times (K\chi) \times (K\chi)$. This costs $O(M^2K^2\chi^2)$ per site, or $O(NM^2K^2\chi^2)$ total. TT-rounding then reduces bond dimensions from $K\chi$ to χ via SVD, costing $O(NM(K\chi)^3) = O(NMK^3\chi^3)$ in the worst case. For one time step of an Euler or Runge-Kutta scheme, the total cost is $O(NMK^3\chi^3)$.

For fixed K, M, χ, the explicit method's $O(M^N)$ dependence on N is exponential, while the TT method's $O(N)$ dependence is linear. This is the claimed exponential-to-linear separation.

The separation is conditional on the bond dimension χ remaining bounded as N grows. For CME solutions arising from biologically structured networks with local interactions, empirical and theoretical evidence suggests that χ grows at most polynomially in N and the desired accuracy [11, 10], though proving sharp bounds on χ for general reaction networks remains open. □

B Implementation Details

B.1 Software Architecture

TN-Check is implemented in Python with NumPy and SciPy for numerical linear algebra. The implementation is structured into the following modules:

`generator/` Constructs the MPO rate matrix from a reaction network specification. Supports mass-action kinetics with configurable rate constants and per-species copy-number bounds.

`solver/` Time-stepping routines for the CME in TT format. Includes Euler and Runge-Kutta integrators with adaptive TT-rounding after each step. Also provides steady-state solvers via power iteration.

`verifier/` The CSL model-checking engine. Implements satisfaction-set computation for all CSL operators. Atomic propositions are constructed as rank-1 MPS vectors (axis-aligned) or bounded-rank MPS vectors (linear predicates). The bounded-until operator uses matrix exponential computation (dense for small systems, Krylov-based for larger ones).

`certificate/` The independent certificate verification module. Performs the 11 checks described in Section 5.4. Outputs a structured verdict with pass/fail status for each check and the total certified error.

`evaluator/` Utility functions for extracting probabilities, means, and other statistics from MPS vectors. Includes efficient contraction routines for computing inner products and expectation values.

B.2 MPS Operations

Key MPS operations and their implementations:

- **Inner product** $\langle \mathbf{u}, \mathbf{v} \rangle$: Left-to-right contraction of the transfer matrices $\sum_{j_k} \bar{U}_k(j_k) \otimes V_k(j_k)$, cost $O(NM\chi^4)$.

- **MPO-MPS product** $Q\mathbf{v}$: Site-by-site contraction producing an MPS with bond dimensions $r_k \cdot s_k$, followed by TT-rounding.
- **Element-wise product** (Hadamard): Kronecker product of cores at each site, producing bond dimensions $r_k \cdot r'_k$, followed by TT-rounding.
- **TT-rounding**: Left-to-right QR sweep followed by right-to-left SVD sweep with truncation. The truncation threshold ε controls the trade-off between bond dimension and accuracy.
- **Clamping**: For small systems, entries are extracted, clamped to $[0, \infty)$, and re-encoded as MPS. For large systems, the alternating-projection scheme of Algorithm 1 is used.

B.3 Numerical Considerations

- **Normalization:** After each TT-rounding step, the MPS is renormalized so that $\|\mathbf{p}\|_1 = 1$ (for probability vectors). The normalization constant is tracked for error accounting.
- **Gauge fixing:** MPS cores are maintained in left-canonical or right-canonical form to ensure numerical stability of the SVD-based rounding.
- **Bond dimension control:** A maximum bond dimension $\chi_{\max}$ is enforced throughout. Adaptive truncation increases $\chi_{\max}$ if the truncation error exceeds a user-specified threshold.

C Full Experimental Data

C.1 Birth-Death Process Details

Model specification:

- Species: X (single species, $N = 1$).
- Reactions: $\emptyset \xrightarrow{\lambda} X$ (birth, $\lambda = 10$), $X \xrightarrow{\mu} \emptyset$ (death, $\mu = 1$ per molecule).
- Copy-number bound: $M = 30$ (`max_copy` $= 30$).
- Stationary distribution: Poisson(λ/μ) = Poisson(10) truncated to $\{0, \ldots, 29\}$.

Results:

- Mean (exact Poisson): 9.999998.
- Mean (MPS): 9.999998.
- L1 error: 0.0 (exact agreement to machine precision).
- Maximum bond dimension: 1 (single site, trivial bonds).

The slight deviation of the mean from 10.0 is due to the truncation of the Poisson distribution at $M = 30$: the tail probability $\Pr(X \geq 30) \approx 7.8 \times 10^{-7}$ when $X \sim$ Poisson(10) is redistributed, causing a negligible shift.

C.2 Compression Accuracy Details

The gene expression model consists of two species (mRNA X_1 and protein X_2) with reactions:

$$\emptyset \xrightarrow{k_1} X_1, \qquad X_1 \xrightarrow{\delta_1} \emptyset,$$
$$X_1 \xrightarrow{k_2} X_1 + X_2, \qquad X_2 \xrightarrow{\delta_2} \emptyset.$$

With per-species bounds $M_1 = 20$ (mRNA) and $M_2 = 50$ (protein), the state space has $|\mathcal{S}| = 20 \times 50 = 1{,}000$ states.

The exact stationary distribution was computed via dense null-space computation (`scipy.linalg.null_space` on the 1000×1000 rate matrix) and then compressed to MPS form with bond dimensions $\chi \in \{1, 2, 5, 10\}$.

Table 11: Full compression accuracy data for gene expression model.

χ	L1 error	MPS params	Compression ratio
1	2.35×10^{-1}	70	14.3×
2	4.71×10^{-2}	280	3.6×
5	4.88×10^{-5}	1,750	0.57×
10	1.62×10^{-12}	7,000	0.14×

For this small system ($N = 2$), the MPS representation is larger than the dense vector for $\chi \geq 5$. The compression advantage of MPS emerges for larger N: at $N = 10$ with $M = 10$, the dense vector has 10^{10} entries while the MPS with $\chi = 10$ has $10 \times 10 \times 10^2 = 10{,}000$ entries per site, totaling $\sim 10^5$ parameters.

C.3 Clamping Bound Verification Details

Experimental protocol:

1. Generate a random non-negative vector $\mathbf{p} \in \mathbb{R}^{d_1 \times d_2}$ with $d_1, d_2 \in \{4, 5, 6, 7, 8\}$ (chosen uniformly at random).
2. Normalize: $\mathbf{p} \leftarrow \mathbf{p}/\|\mathbf{p}\|_1$.
3. Represent as a 2-site MPS (exact, bond dimension $\leq \min(d_1, d_2)$).
4. Truncate aggressively to bond dimension $\chi \in \{1, 2, 3\}$ (chosen uniformly at random), obtaining $\tilde{\mathbf{p}}$.
5. Compute $\varepsilon_{\text{trunc}} = \|\mathbf{p} - \tilde{\mathbf{p}}\|_1$.
6. Clamp: $\hat{\mathbf{p}} = \max(\tilde{\mathbf{p}}, 0)$.
7. Compute $\varepsilon_{\text{clamp}} = \|\mathbf{p} - \hat{\mathbf{p}}\|_1$.
8. Check: $\varepsilon_{\text{clamp}} \leq \varepsilon_{\text{trunc}} + 10^{-10}$.

All 10 trials passed the check, confirming Proposition 6. The tolerance 10^{-10} accounts for floating-point arithmetic.

Error propagation was additionally verified on the birth-death model with bond dimensions $\chi \in \{1, 2, 3, 4, 5\}$: in all cases, Proposition 6 was satisfied, confirming that clamping consistently reduces L1 error across different bond dimensions.

C.4 Scaling Data Details

The scaling experiments (Table 3) use N independent birth-death species, each with $M = 10$ states and reactions $\emptyset \xrightarrow{\lambda} X_i$, $X_i \xrightarrow{\mu} \emptyset$. The MPO is compiled from the sum-of-Kronecker-products decomposition (2), with each reaction contributing one term.

The MPO parameter count is computed as $\sum_{k=1}^{N} M^2 \cdot s_{k-1} \cdot s_k$ where s_k are the MPO bond dimensions. The compression ratio is M^N/MPO params, representing the ratio of the dense state space to the MPO parameter count.

Note that the compression ratio in Table 3 compares the MPO parameter count to the *dense state space* M^N, not to the sparse rate matrix. The sparse rate matrix has $O(KM^N)$ nonzero entries, so the compression ratio relative to sparse storage is smaller by a factor of K. Nevertheless, the exponential growth of the compression ratio with N holds in either case.

C.5 Multi-Species CSL Details

The 3-species cascade model:

$$\begin{aligned} &\emptyset \xrightarrow{k_0} X_0, \qquad X_0 \xrightarrow{\delta_0} \emptyset, \\ &X_0 \xrightarrow{k_1} X_0 + X_1, \qquad X_1 \xrightarrow{\delta_1} \emptyset, \\ &X_1 \xrightarrow{k_2} X_1 + X_2, \qquad X_2 \xrightarrow{\delta_2} \emptyset. \end{aligned}$$

With per-species bound $M = 8$, the state space is $\{0, \ldots, 7\}^3$ with $|\mathcal{S}| = 512$.

Dense ground-truth computation: the exact probability $\Pr(\text{reach } X_2 \geq 3 \text{ within } t = 2) = 0.006$ was computed via dense matrix exponential on the 512×512 rate matrix after making non-goal states absorbing.

The TT-based computation returned a probability interval of $[0, 1]$ with L1 error bound 1.98. Since $1.98 > 0.006$, the error bound exceeds the true probability, and the checker correctly returns "indeterminate."

C.6 Certificate Verification Details

Configuration 1: 2 species, 10×10. Physical dimensions $d_1 = d_2 = 10$, state space size = 100. Certificate checks:

- Checks C1–C11: all pass.
- Verdict: SOUND.
- Total certified error: 8.50×10^{-3}.

Configuration 2: 2 species, end-to-end. Full pipeline including time evolution and CSL evaluation. Certificate checks:

- Checks C1–C11: all pass.
- Verdict: SOUND.
- Total certified error: 2.38×10^{-2}.

The higher total error in Configuration 2 reflects the accumulation of errors across the full pipeline (time stepping + TT rounding + CSL evaluation), compared to Configuration 1 which tests only the static certificate on a pre-computed distribution.

C.7 Proof of Theorem 11 (Deficiency-Zero Compressibility)

Proof. We prove each part:

Part (i): Product-form stationary distribution. By the Deficiency-Zero Theorem (Feinberg [6]), a weakly reversible mass-action system with deficiency $\delta = 0$ has a unique positive equilibrium in each stoichiometric compatibility class, and this equilibrium is complex-balanced. Anderson, Craciun, and Kurtz [1] extended this to the stochastic setting: for a complex-balanced mass-action system, the stationary distribution of the corresponding CME is product-form: $\pi(\mathbf{x}) = \prod_{i=1}^{N} \frac{c_i^{x_i}}{x_i!} e^{-c_i}$ where c_i is the deterministic equilibrium concentration of species i. Each marginal $\pi_i(x_i) = \frac{c_i^{x_i}}{x_i!} e^{-c_i}$ is Poisson.

Part (ii): TT rank of π. A product-form distribution $\pi = \bigotimes_{i=1}^{N} \pi_i$ is by definition a rank-1 tensor in the tensor-train decomposition: each core $G_k(j_k) = \pi_k(j_k)$ is a 1×1 matrix, giving bond dimensions $r_0 = r_1 = \cdots = r_N = 1$.

Part (iii): Transient TT rank bound. Write $\mathbf{p}(t) = \pi + \mathbf{e}(t)$ where $\mathbf{e}(t) = e^{Qt}(\mathbf{p}(0) - \pi)$ is the deviation from equilibrium. By the spectral decomposition of Q, $\|\mathbf{e}(t)\|_1 \leq C_0 e^{-\gamma t}$ where $\gamma = \mathrm{Re}(-\lambda_2)$ is the spectral gap (second eigenvalue of Q) and $C_0 = \|\mathbf{p}(0) - \pi\|_1$. At TT-SVD tolerance ε, the rank required to represent $\mathbf{p}(t)$ is the rank of π (which is 1) plus the rank needed for the error term. When $C_0 e^{-\gamma t} < \varepsilon$, the error term is below truncation tolerance and $\chi = 1$ suffices. For intermediate times, $\chi \leq 1 + C_{\mathrm{mix}} e^{-\gamma t}$ where C_{mix} depends on the Schmidt spectrum of the initial deviation $\mathbf{p}(0) - \pi$. □

C.8 Proof of Theorem 13 (Unified Error Propagation)

Proof. We track the L1 error through K time steps. Let $\mathbf{p}_k$ denote the exact probability vector at time $t_k = k\Delta t$ and $\tilde{\mathbf{p}}_k$ the computed approximation.

Step 1: Error decomposition per step. At step k, three error sources contribute:

(a) *TT-SVD truncation:* After applying the operator and rounding, $\|\mathbf{q}_k - \mathrm{round}(\mathbf{q}_k)\|_1 \leq \varepsilon_{\mathrm{TT}}^{(k)}$.

(b) *TDVP projection:* The tangent-space projection introduces a Frobenius residual $r_k = \|\partial_t \psi_k - Q\psi_k\|_F$. By Proposition 14, this converts to L1 error $\leq \sqrt{2K_{\mathrm{eff}}} \cdot r_k \Delta t$.

(c) *Clamping:* By Proposition 6, clamping does not increase L1 error, so it contributes 0 to the bound.

Step 2: Error propagation across steps. By Theorem 1, the Metzler semigroup satisfies $C(t) = 1$ for all $t \geq 0$. Therefore, the error from step k propagates to step K without amplification: $\|e^{Q(T-t_k)}\mathbf{e}_k\|_1 \leq \|\mathbf{e}_k\|_1$.

Step 3: Total error. Summing over all sources:

$$\|\mathbf{p}(T) - \tilde{\mathbf{p}}(T)\|_1 \leq \varepsilon_{\mathrm{FSP}} + \sum_{k=1}^{K} \varepsilon_{\mathrm{TT}}^{(k)} + \sqrt{2K_{\mathrm{eff}}} \sum_{k=1}^{K} r_k \Delta t. \tag{15}$$

The FSP truncation contributes $\varepsilon_{\mathrm{FSP}}$ at initialization (before time stepping begins). The key property is *linear* accumulation: K steps contribute $K \cdot \varepsilon_{\text{per-step}}$, not $C^K \cdot \varepsilon_{\text{per-step}}$, because $C(t) = 1$ for Metzler generators.

Step 4: Verdict certification. For a CSL property $\mathcal{P}_{\geq p}[\phi]$ with computed probability $\hat{p}$, define the *verdict margin* $m = |\hat{p} - p|$. If $m > \varepsilon_{\mathrm{total}}$, the true probability p^* satisfies $|p^* - \hat{p}| \leq \varepsilon_{\mathrm{total}} < m$, so p^* is on the same side of the threshold as $\hat{p}$, and the verdict is certified sound. If $m \leq \varepsilon_{\mathrm{total}}$, the three-valued semantics returns INDETERMINATE, which is always sound. □